Introduction to SQL

Practice Exercises

3.1 Write the following queries in SQL, using the university schema. (We suggest you actually run these queries on a database, using the sample data that we provide on the web site of the book, db-book.com. Instructions for setting up a database, and loading sample data, are provided on the above web site.)

    a. Find the titles of courses in the Comp. Sci. department that have 3 credits.

    b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.

    c. Find the highest salary of any instructor.

    d. Find all instructors earning the highest salary (there may be more than one with the same salary).

    e. Find the enrollment of each section that was offered in Fall 2017.

    f. Find the maximum enrollment, across all sections, in Fall 2017.

    g. Find the sections that had the maximum enrollment in Fall 2017.

Answer:

    a. Find the titles of courses in the Comp. Sci. department that have 3 credits.

> **select** *title*
> **from** *course*
> **where** *dept_name* = 'Comp. Sci.' **and** *credits* = 3

    b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.
This query can be answered in several different ways. One way is as follows.

> **select** **distinct** *takes.ID*
> **from** *takes, instructor, teaches*
> **where** *takes.course_id* = *teaches.course_id* **and**
> *takes.sec_id* = *teaches.sec_id* **and**
> *takes.semester* = *teaches.semester* **and**
> *takes.year* = *teaches.year* **and**
> *teaches.id* = *instructor.id* **and**
> *instructor.name* = 'Einstein'

c.   Find the highest salary of any instructor.

> **select max**(*salary*)
> **from** *instructor*

d.   Find all instructors earning the highest salary (there may be more than one with the same salary).

> **select** *ID*, *name*
> **from** *instructor*
> **where** *salary* = (**select max**(*salary*) **from** *instructor*)

e.   Find the enrollment of each section that was offered in Fall 2017.

> **select** *course_id*, *sec_id*,
> (**select count**(*ID*)
> **from** *takes*
> **where** *takes.year* = *section.year*
> **and** *takes.semester* = *section.semester*
> **and** *takes.course_id* = *section.course_id*
> **and** *takes.sec_id* = *section.sec_id*)
> **as** *enrollment*
> **from** *section*
> **where** *semester* = 'Fall'
> **and** *year* = 2017

Note that if the result of the subquery is empty, the aggregate function **count** returns a value of 0.

One way of writing the query might appear to be:

```
select   takes.course_id, takes.sec_id, count(ID)
from     section, takes
where    takes.course_id = section.course_id
         and takes.sec_id = section.sec_id
         and takes.semester = section.semester
         and takes.year = section.year
         and takes.semester = 'Fall'
         and takes.year = 2017
group by takes.course_id, takes.sec_id
```

But note that if a section does not have any students taking it, it would not appear in the result. One way of ensuring such a section appears with a count of 0 is to use the **outer join** operation, covered in Chapter 4.

f. Find the maximum enrollment, across all sections, in Fall 2017.
One way of writing this query is as follows:

```
select max(enrollment)
from   (select  count(ID) as enrollment
        from    section, takes
        where   takes.year = section.year
                and takes.semester = section.semester
                and takes.course_id = section.course_id
                and takes.sec_id = section.sec_id
                and takes.semester = 'Fall'
                and takes.year = 2017
                group by takes.course_id, takes.sec_id)
```

As an alternative to using a nested subquery in the **from** clause, it is possible to use a **with** clause, as illustrated in the answer to the next part of this question.
   A subtle issue in the above query is that if no section had any enrollment, the answer would be empty, not 0. We can use the alternative using a subquery, from the previous part of this question, to ensure the count is 0 in this case.

g. Find the sections that had the maximum enrollment in Fall 2017.
The following answer uses a **with** clause, simplifying the query.

```
with  sec_enrollment as (
      select   takes.course_id, takes.sec_id, count(ID) as enrollment
      from     section, takes
      where    takes.year = section.year
               and takes.semester = section.semester
               and takes.course_id = section.course_id
               and takes.sec_id = section.sec_id
               and takes.semester = 'Fall'
               and takes.year = 2017
      group by takes.course_id, takes.sec_id)
select   course_id, sec_id
from     sec_enrollment
where    enrollment = (select max(enrollment) from sec_enrollment)
```

It is also possible to write the query without the **with** clause, but the sub-query to find enrollment would get repeated twice in the query.

While not incorrect to add **distinct** in the **count**, it is not necessary in light of the primary key constraint on *takes*.

3.2   Suppose you are given a relation *grade_points*(*grade*, *points*) that provides a conversion from letter grades in the *takes* relation to numeric scores; for example, an "A" grade could be specified to correspond to 4 points, an "A−" to 3.7 points, a "B+" to 3.3 points, a "B" to 3 points, and so on. The grade points earned by a student for a course offering (section) is defined as the number of credits for the course multiplied by the numeric points for the grade that the student received.

Given the preceding relation, and our university schema, write each of the following queries in SQL. You may assume for simplicity that no *takes* tuple has the *null* value for *grade*.

   a.   Find the total grade points earned by the student with ID '12345', across all courses taken by the student.

   b.   Find the grade point average (*GPA*) for the above student, that is, the total grade points divided by the total credits for the associated courses.

   c.   Find the ID and the grade-point average of each student.

   d.   Now reconsider your answers to the earlier parts of this exercise under the assumption that some grades might be null. Explain whether your solutions still work and, if not, provide versions that handle nulls properly.

Answer:

   a.   Find the total grade-points earned by the student with ID '12345', across all courses taken by the student.

> **select sum**(*credits* * *points*)
> **from** *takes, course, grade_points*
> **where** *takes.grade = grade_points.grade*
>   **and** *takes.course_id = course.course_id*
>   **and** *ID* = '12345'

In the above query, a student who has not taken any course would not have any tuples, whereas we would expect to get 0 as the answer. One way of fixing this problem is to use the **outer join** operation, which we study later in Chapter 4. Another way to ensure that we get 0 as the answer is via the following query:

> (**select**   **sum**(*credits* * *points*)
> **from**   *takes, course, grade_points*
> **where**   *takes.grade = grade_points.grade*
>   **and** *takes.course_id = course.course_id*
>   **and** *ID*= '12345')
> **union**
> (**select**   0
> **from**   *student*
> **where**   *ID* = '12345' **and**
>   **not exists** (**select** * **from** *takes* **where** *ID* = '12345'))

b. Find the grade point average (*GPA*) for the above student, that is, the total grade-points divided by the total credits for the associated courses.

> **select**   **sum**(*credits* * *points*)/**sum**(*credits*) **as** *GPA*
> **from**   *takes, course, grade_points*
> **where**   *takes.grade = grade_points.grade*
>   **and** *takes.course_id = course.course_id*
>   **and** *ID*= '12345'

As before, a student who has not taken any course would not appear in the above result; we can ensure that such a student appears in the result by using the modified query from the previous part of this question. However, an additional issue in this case is that the sum of credits would also be 0, resulting in a divide-by-zero condition. In fact, the only meaningful way of defining the *GPA* in this case is to define it as *null*. We can ensure that such a student appears in the result with a null *GPA* by adding the following **union** clause to the above query.

> **union**
> (**select** *null* **as** *GPA*
> **from**   *student*
> **where** *ID* = '12345' **and**
>   **not exists** (**select** * **from** *takes* **where** *ID* = '12345'))

c. Find the ID and the grade-point average of each student.

> **select**   *ID*, **sum**(*credits * points*)/**sum**(*credits*) **as** *GPA*
> **from**   *takes, course, grade_points*
> **where**   *takes.grade = grade_points.grade*
>   **and** *takes.course_id = course.course_id*
> **group by** *ID*

Again, to handle students who have not taken any course, we would have to add the following **union** clause:

> **union**
> (**select** *ID*, *null* **as** *GPA*
> **from**   *student*
> **where** **not exists** (**select * from** *takes* **where** *takes.ID = student.ID*))

d. Now reconsider your answers to the earlier parts of this exercise under the assumption that some grades might be null. Explain whether your solutions still work and, if not, provide versions that handle nulls properly. The queries listed above all include a test of equality on *grade* between *grade_points* and *takes*. Thus, for any *takes* tuple with a *null* grade, that student's course would be eliminated from the rest of the computation of the result. As a result, the credits of such courses would be eliminated also, and thus the queries would return the correct answer even if some grades are null.

3.3 Write the following inserts, deletes, or updates in SQL, using the university schema.

a. Increase the salary of each instructor in the Comp. Sci. department by 10%.

b. Delete all courses that have never been offered (i.e., do not occur in the *section* relation).

c. Insert every student whose *tot_cred* attribute is greater than 100 as an instructor in the same department, with a salary of $10,000.

a. Increase the salary of each instructor in the Comp. Sci. department by 10%.

> **update** *instructor*
> **set**   *salary = salary * 1.10*
> **where** *dept_name* = 'Comp. Sci.'

b. Delete all courses that have never been offered (that is, do not occur in the *section* relation).

person (*driver_id*, *name*, *address*)
car (*license_plate*, *model*, *year*)
accident (*report_number*, *year*, *location*)
owns (*driver_id*, *license_plate*)
participated (*report_number*, *license_plate*, *driver_id*, *damage_amount*)

Figure 3.17   Insurance database

> **delete** **from** *course*
> **where** *course_id* **not in**
>   (**select** *course_id* **from** *section*)

c. Insert every student whose *tot_cred* attribute is greater than 100 as an instructor in the same department, with a salary of $10,000.

> **insert into** *instructor*
> **select** *ID*, *name*, *dept_name*, 10000
> **from** *student*
> **where** *tot_cred* > 100

3.4 Consider the insurance database of Figure 3.17, where the primary keys are underlined. Construct the following SQL queries for this relational database.

a. Find the total number of people who owned cars that were involved in accidents in 2017.

b. Delete all year-2010 cars belonging to the person whose ID is '12345'.

Answer:

a. Find the total number of people who owned cars that were involved in accidents in 2017.
Note: This is not the same as the total number of accidents in 2017. We must count people with several accidents only once. Furthermore, note that the question asks for owners, and it might be that the owner of the car was not the driver actually involved in the accident.

> **select** count (**distinct** *person.driver_id*)
> **from** *accident*, *participated*, *person*, *owns*
> **where** *accident.report_number* = *participated.report_number*
>   **and** *owns.driver_id* = *person.driver_id*
>   **and** *owns.license_plate* = *participated.license_plate*
>   **and** *year* = 2017

b.   Delete all year-2010 cars belonging to the person whose ID is '12345'.

> **delete** *car*
> **where** *year* = 2010 **and** *license_plate* **in**
>   (**select** *license_plate*
>   **from** *owns o*
>   **where** *o.driver_id* = '12345')

Note: The *owns*, *accident* and *participated* records associated with the deleted cars still exist.

3.5   Suppose that we have a relation *marks*(*ID*, *score*) and we wish to assign grades to students based on the score as follows: grade *F* if *score* < 40, grade *C* if 40 ≤ *score* < 60, grade *B* if 60 ≤ *score* < 80, and grade *A* if 80 ≤ *score*. Write SQL queries to do the following:

a.   Display the grade for each student, based on the *marks* relation.

b.   Find the number of students with each grade.

a.   Display the grade for each student, based on the *marks* relation.

> **select** *ID*,
>   **case**
>   **when** *score* < 40 **then** 'F'
>   **when** *score* < 60 **then** 'C'
>   **when** *score* < 80 **then** 'B'
>   **else** 'A'
>   **end**
> **from** *marks*

b.   Find the number of students with each grade.

```
with     grades as
(
select   ID,
         case
              when score < 40 then 'F'
              when score < 60 then 'C'
              when score < 80 then 'B'
              else 'A'
         end as grade
from     marks
)
select   grade, count(ID)
from     grades
group by grade
```

As an alternative, the **with** clause can be removed, and instead the definition of *grades* can be made a subquery of the main query.

3.6 The SQL **like** operator is case sensitive (in most systems), but the lower() function on strings can be used to perform case-insensitive matching. To show how, write a query that finds departments whose names contain the string "sci" as a substring, regardless of the case.

Answer:

```
select  dept_name
from    department
where   lower(dept_name) like '%sci%'
```

3.7 Consider the SQL query

```
select p.a1
from p, r1, r2
where p.a1 = r1.a1 or p.a1 = r2.a1
```

Under what conditions does the preceding query select values of $p.a1$ that are either in $r1$ or in $r2$? Examine carefully the cases where either $r1$ or $r2$ may be empty.

Answer:
The query selects those values of $p.a1$ that are equal to some value of $r1.a1$ or $r2.a1$ if and only if both $r1$ and $r2$ are non-empty. If one or both of $r1$ and $r2$ are empty, the Cartesian product of $p, r1$ and $r2$ is empty, hence the result of the query is empty. If p itself is empty, the result is empty.

3.8 Consider the bank database of Figure 3.18, where the primary keys are underlined. Construct the following SQL queries for this relational database.

---

branch(*branch_name*, *branch_city*, *assets*)
customer (*ID*, *customer_name*, *customer_street*, *customer_city*)
loan (*loan_number*, *branch_name*, *amount*)
borrower (*ID*, *loan_number*)
account (*account_number*, *branch_name*, *balance*)
depositor (*ID*, *account_number*)

Figure 3.18   Banking database.

a.   Find the ID of each customer of the bank who has an account but not a loan.

b.   Find the ID of each customer who lives on the same street and in the same city as customer '12345'.

c.   Find the name of each branch that has at least one customer who has an account in the bank and who lives in "Harrison".

Answer:

a.   Find the ID of each customer of the bank who has an account but not a loan.

$$
\begin{aligned}
&(\textbf{select}\ \ ID \\
&\ \textbf{from}\ \ \ \ \ \ depositor) \\
&\ \textbf{except} \\
&(\textbf{select}\ \ ID \\
&\ \textbf{from}\ \ \ \ \ \ borrower)
\end{aligned}
$$

b.   Find the ID of each customer who lives on the same street and in the same city as customer '12345'.

$$
\begin{aligned}
&\textbf{select}\ \ \ F.ID \\
&\textbf{from}\ \ \ \ \ customer\ \textbf{as}\ F,\ customer\ \textbf{as}\ S \\
&\textbf{where}\ \ \ F.customer\_street = S.customer\_street \\
&\qquad\quad\ \ \textbf{and}\ F.customer\_city = S.customer\_city \\
&\qquad\quad\ \ \textbf{and}\ S.customer\_id\ = \text{'}12345\text{'}
\end{aligned}
$$

c.   Find the name of each branch that has at least one customer who has an account in the bank and who lives in "Harrison".

> **select** **distinct** *branch_name*
> **from** *account, depositor, customer*
> **where** *customer.id = depositor.id*
> **and** *depositor.account_number = account.account_number*
> **and** *customer_city* = 'Harrison'

3.9 Consider the relational database of Figure 3.19, where the primary keys are underlined. Give an expression in SQL for each of the following queries.

a. Find the ID, name, and city of residence of each employee who works for "First Bank Corporation".

b. Find the ID, name, and city of residence of each employee who works for "First Bank Corporation" and earns more than $10000.

c. Find the ID of each employee who does not work for "First Bank Corporation".

d. Find the ID of each employee who earns more than every employee of "Small Bank Corporation".

e. Assume that companies may be located in several cities. Find the name of each company that is located in every city in which "Small Bank Corporation" is located.

f. Find the name of the company that has the most employees (or companies, in the case where there is a tie for the most).

g. Find the name of each company whose employees earn a higher salary, on average, than the average salary at "First Bank Corporation".

Answer:

a. Find the ID, name, and city of residence of each employee who works for "First Bank Corporation".

> *employee* (*ID*, *person_name*, *street*, *city*)
> *works* (*ID*, *company_name*, *salary*)
> *company* (*company_name*, *city*)
> *manages* (*ID*, *manager_id*)

Figure 3.19 Employee database.

> **select** *e.ID, e.person_name, city*
> **from** *employee* **as** *e*, *works* **as** *w*
> **where** *w.company_name* = 'First Bank Corporation' **and**
>   *w.ID* = *e.ID*

b. Find the ID, name, and city of residence of each employee who works for "First Bank Corporation" and earns more than $10000.

select *
from *employee*
where *ID* **in**
   (**select** *ID*
    from *works*
    where *company_name* = 'First Bank Corporation' **and** *salary > 10000*)

This could be written also in the style of the answer to part *a*.

c. Find the ID of each employee who does not work for "First Bank Corporation".

> **select** *ID*
> **from** *works*
> **where** *company_name* <> 'First Bank Corporation'

If one allows people to appear in *employee* without appearing also in *works*, the solution is slightly more complicated. An outer join as discussed in Chapter 4 could be used as well.

> **select** *ID*
> **from** *employee*
> **where** *ID* **not in**
>   (**select** *ID*
>   **from** *works*
>   **where** *company_name* = 'First Bank Corporation')

d. Find the ID of each employee who earns more than every employee of "Small Bank Corporation".

> **select** *ID*
> **from** *works*
> **where** *salary* > **all**
>   (**select** *salary*
>   **from** *works*
>   **where** *company_name* = 'Small Bank Corporation')

If people may work for several companies and we wish to consider the *total* earnings of each person, the problem is more complex. But note that the

fact that ID is the primary key for *works* implies that this cannot be the case.

e. Assume that companies may be located in several cities. Find the name of each company that is located in every city in which "Small Bank Corporation" is located.

> **select** *S.company_name*
> **from** *company* **as** *S*
> **where not exists** ((**select** *city*
>   **from** *company*
>   **where** *company_name* = 'Small Bank Corporation')
>   **except**
>   (**select** *city*
>   **from** *company* **as** *T*
>   **where** *S.company_name* = *T.company_name*))

f. Find the name of the company that has the most employees (or companies, in the case where there is a tie for the most).

> **select** *company_name*
> **from** *works*
> **group by** *company_name*
> **having count** (**distinct** *ID*) >= **all**
>   (**select count** (**distinct** *ID*)
>   **from** *works*
>   **group by** *company_name*)

g. Find the name of each company whose employees earn a higher salary, on average, than the average salary at "First Bank Corporation".

> **select** *company_name*
> **from** *works*
> **group by** *company_name*
> **having avg** (*salary*) > (**select avg** (*salary*)
>   **from** *works*
>   **where** *company_name* = 'First Bank Corporation')

3.10 Consider the relational database of Figure 3.19. Give an expression in SQL for each of the following:

a. Modify the database so that the employee whose ID is '12345' now lives in "Newtown".

b. Give each manager of "First Bank Corporation" a 10 percent raise unless the salary becomes greater than $100000; in such cases, give only a 3 percent raise.

Answer:

a.   Modify the database so that the employee whose ID is '12345' now lives in "Newtown".

> **update** *employee*
> **set** *city* = 'Newtown'
> **where** *ID* = '12345'

b.   Give each manager of "First Bank Corporation" a 10 percent raise unless the salary becomes greater than $100000; in such cases, give only a 3 percent raise.

> **update** *works T*
> **set** *T.salary* = *T.salary* * 1.03
> **where** *T*.*ID* **in** (**select** *manager_id*
>   **from** *manages*)
>   **and** *T.salary* * 1.1 > 100000
>   **and** *T.company_name* = 'First Bank Corporation'

> **update** *works T*
> **set** *T.salary* = *T.salary* * 1.1
> **where** *T*.*ID* **in** (**select** *manager_id*
>   **from** *manages*)
>   **and** *T.salary* * 1.1 <= 100000
>   **and** *T.company_name* = 'First Bank Corporation'

The above updates would give different results if executed in the opposite order. We give below a safer solution using the **case** statement.

> **update** *works T*
> **set** *T.salary* = *T.salary* $*$
>   (**case**
>   **when** (*T.salary* $*$ 1.1 > 100000) **then** 1.03
>   **else** 1.1
>   **end**)
> **where** *T.ID* **in** (**select** *manager_id*
>   **from** *manages*) **and**
>   *T.company_name* = 'First Bank Corporation'